

# Database Systems

Tutorial Week 5

# Objectives

- I. Relational algebra (RA) review
- II. Relational algebra and SQL statements

# Relational Algebra

- Theory behind SQL
- Gives a plan for evaluating a query
- Exploiting “equivalencies” of relational operators can lead to faster queries
  - E.g. joins can be expensive, but we can “push” selections and projections ahead of joins so you’re joining on a smaller table
- Input to an operation?
  - Table(s) of a relation with rows and columns
- Output?
  - A table of a relation with rows and columns

# Fundamental Operations

- 5 basic operators of Relational Algebra
- Can form other compound operations
- What are they?
  - Removal operators:
    - Projection ( $\pi$ )
    - Selection ( $\sigma$ )
  - Set operators:
    - Union ( $\cup$ )
    - Set difference ( $-$ )
    - Cross product ( $\times$ )

# Removal Operators

- Remove components from a table of a relation
- Projection ( $\pi$ )
  - Removes columns
  - $\pi_{A1, A2, \dots, An}(R)$  where  $R$  = relation and  $A1, A2, \dots, An$  are attributes that are “projected”
  - Creates a new relation with a subset of attributes
  - All the tuples are included in the new relation, but only the attributes  $A1, A2, \dots, An$  are kept
  - SQL: read  $\pi$  as `SELECT`
- Selection ( $\sigma$ )
  - Removes rows
  - $\sigma_C(R)$  where  $R$  = relation and  $C$  = condition used to filter rows
  - Creates a new relation with rows where  $C$  is true
  - SQL: read  $\sigma$  as `WHERE`

# Example Time

Consider the following Person table:

FirstName	LastName	Phone	Email
Jon	Snow	0551-999-210	knowsnothing@hotmail.com
Daenerys	Targaryen	0569-988-112	bendtheknee@gmail.com
Jamie	Lannister	0531-987-654	handsfree@gmail.com
Night	King	0566-123-456	killerstare@gmail.com

What will  $\pi_{(\text{FirstName}, \text{LastName})}(\text{Person})$  result in?

# Example Time

What will  $\pi_{(\text{FirstName}, \text{LastName})}(\text{Person})$  result in?

<b>FirstName</b>	<b>LastName</b>
Jon	Snow
Daenerys	Targaryen
Jamie	Lannister
Night	King

# Example Time

For the same original Person table, what will  $\sigma_{\text{FirstName} = \text{'Jon'} \vee \text{LastName} = \text{'King'}}(\text{Person})$  result in?

And:  $\wedge$

Or:  $\vee$

FirstName	LastName	Phone	Email
Jon	Snow	0551-999-210	knowsnothing@hotmail.com
Daenerys	Targaryen	0569-988-112	bendtheknee@gmail.com
Jamie	Lannister	0531-987-654	handsfree@gmail.com
Night	King	0566-123-456	killerstare@gmail.com



# Example Time

For the same original Person table, what will  $\sigma_{\text{FirstName} = \text{'Jon'} \vee \text{LastName} = \text{'King'}}(\text{Person})$  result in?

And:  $\wedge$

Or:  $\vee$

FirstName	LastName	Phone	Email
Jon	Snow	0551-999-210	knowsnothing@hotmail.com
Night	King	0566-123-456	killerstare@gmail.com

# Example Time

Can combine the two operations in one expression:

$\pi_{\text{FirstName, LastName}}(\sigma_{\text{FirstName} = \text{'Jon'} \vee \text{LastName} = \text{'King'}}(\text{Person}))$

FirstName	LastName
Jon	Snow
Night	King

# Set Operators

- Operate on 2 relations
- Union and Difference operations take two input relations, which must be **union-compatible**:
  - They must have the same attribute names in the same order
  - Corresponding attributes must have the same datatype
- Union ( $\cup$ )
  - $R \cup S$  where R and S are relations
  - Result: every row which is either in R or S
- Set difference ( $-$ )
  - $R - S$  where R and S are relations
  - Result: every row which is in R but not in S
  - Set-difference is **not** symmetrical:  $R - S \neq S - R$ !!!!!!

# Example Time

Consider the following tables, GoodGuys and BadGuys:

GoodGuys

FirstName	LastName
Jon	Snow
Daenerys	Targaryen

BadGuys

FirstName	LastName
Cersei	Lannister
Night	King

What will  $\text{GoodGuys} \cup \text{BadGuys}$  result in?

# Example Time

<b>FirstName</b>	<b>LastName</b>
Jon	Snow
Daenerys	Targaryen
Cersei	Lannister
Night	King

# Example Time

Consider the following tables, RandomCombo1 and RandomCombo2:

RandomCombo1

FirstName	LastName
Jon	Snow
Daenerys	Targaryen
Jamie	Lannister
Night	King

RandomCombo2

FirstName	LastName
Night	King
Arya	Stark
Cersei	Lannister
Daenerys	Targaryen

What will RandomCombo1 – RandomCombo2 result in?

# Example Time

<b>FirstName</b>	<b>LastName</b>
Jon	Snow
Jamie	Lannister

# Set Operators

- Cross product ( $\times$ )
  - $R \times S$  where  $R$  and  $S$  are relations
  - Each row of  $R$  pairs with each row of  $S$
  - Resulting schema has all attributes from both relations
  - If some attributes have the same name, rename them by using the renaming operator (will cover later in the tute)



# Example Time

Consider the following tables, Person and Weapon:

Person

FirstName	LastName	Email
Jon	Snow	knowsnothing@hotmail.com
Night	King	killerstare@gmail.com

Weapon

Weapon	Metal
Sword	Valyrian steel
Dagger	Dragon glass

What will Person  $\times$  Weapon result in?

# Example Time

Each row of Person pairs with each row of Weapon

<b>FirstName</b>	<b>LastName</b>	<b>Email</b>	<b>Weapon</b>	<b>Metal</b>
Jon	Snow	knowsnothing@hotmail.com	Sword	Valyrian steel
Jon	Snow	knowsnothing@hotmail.com	Dagger	Dragon glass
Night	King	killerstare@gmail.com	Sword	Valyrian steel
Night	King	killerstare@gmail.com	Dagger	Dragon glass

# Compound Operators

- Useful shorthand
- Can be expressed using basic operators
- What are they?
  - Intersection ( $\cap$ )
  - Natural join ( $\bowtie$ )
  - Condition join / theta join / inner join ( $\bowtie_C$ )

# Compound Operators

- Intersection ( $\cap$ )
  - Also a set operator (but not a basic one lol)
  - Takes two input relations
  - Result: a relation containing all tuples which are present in *both* relations
  - Can be expressed using set differences, so the two input relations need to be union-compatible
  - $R \cap S = R - (R - S)$

# Example Time

Consider the following tables, RandomCombo1 and RandomCombo2:

RandomCombo1

FirstName	LastName
Jon	Snow
Daenerys	Targaryen
Jamie	Lannister
Night	King

RandomCombo2

FirstName	LastName
Night	King
Arya	Stark
Cersei	Lannister
Daenerys	Targaryen

What will  $\text{RandomCombo1} \cap \text{RandomCombo2}$  result in?

# Example Time

<b>FirstName</b>	<b>LastName</b>
Daenerys	Targaryen
Night	King

# Example Time

Also,  $\text{RandomCombo1} \cap \text{RandomCombo2}$

$$= \text{RandomCombo1} - (\text{RandomCombo1} - \text{RandomCombo2})$$

RandomCombo1

FirstName	LastName
Jon	Snow
Daenerys	Targaryen
Jamie	Lannister
Night	King

RandomCombo1 – RandomCombo2

FirstName	LastName
Jon	Snow
Jamie	Lannister

–

=

FirstName	LastName
Daenerys	Targaryen
Night	King

# Compound Operators

- Joins
  - Compound operators involving cross product, selection and (sometimes) projection
- Natural join ( $\bowtie$ )
  - Often just called “join”
  - Takes two input relations e.g. R and S
  - Result: a new relation, pairing each tuple from R and S where the common attributes are equal
  - Can be broken down:
    - Compute  $R \times S$
    - Select rows where attributes that appear in both relations have equal values
    - Project all unique attributes and one copy of the common ones
  - If there are no attributes which have the same name, there's nothing to select
    - If this is the case, what will  
SELECT \* FROM Relation1 NATURAL JOIN Relation2  
result in?
      - The cross product! ;)



# Example Time

Consider the following Person and WeaponOwner tables:

Person

FirstName	LastName	Email
Jon	Snow	knowsnothing@hotmail.com
Daenerys	Targaryen	bendtheknee@gmail.com
Tyrion	Lannister	idrinkandiknow@gmail.com
Night	King	killerstare@gmail.com

WeaponOwner

Weapon	LastName	Metal
Sword	Snow	Valyrian steel
Dagger	Lannister	Dragon glass

# Example Time

Person × WeaponOwner (intermediate result)

FirstName	LastName	Email	Weapon	LastName	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Snow	Valyrian steel
Jon	Snow	knowsnothing@hotmail.com	Dagger	Lannister	Dragon glass
Daenerys	Targaryen	bendtheknee@gmail.com	Sword	Snow	Valyrian steel
Daenerys	Targaryen	bendtheknee@gmail.com	Dagger	Lannister	Dragon glass
Tyrion	Lannister	idrinkandiknow@gmail.com	Sword	Snow	Valyrian steel
Tyrion	Lannister	idrinkandiknow@gmail.com	Dagger	Lannister	Dragon glass
Night	King	killerstare@gmail.com	Sword	Snow	Valyrian steel
Night	King	killerstare@gmail.com	Dagger	Lannister	Dragon glass

# Example Time

Person ⋈ WeaponOwner (natural join)

FirstName	LastName	Email	Weapon	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Valyrian steel
Tyrion	Lannister	idrinkandiknow@gmail.com	Dagger	Dragon glass

# Compound Operators

- Condition join / theta join / inner join ( $\bowtie_C$ )
  - Takes two input relations e.g. R and S
  - $R \bowtie_C S$  joins rows from relations R and S such that the Boolean condition C is true
  - Most commonly, C is of the type  $A = B$  (“equi-join”)
  - Can be expressed using basic operators
    - $R \bowtie_C S = \sigma_C(R \times S)$
  - Condition C often refers to equality of attributes e.g.  $R \bowtie_{r.rid = s.sid} S$

# Example Time

Consider the following tables, Person and WeaponOwner:

Person

FirstName	LastName	Email
Jon	Snow	knowsnothing@hotmail.com
Daenerys	Targaryen	bendtheknee@gmail.com
Tyrion	Lannister	idrinkandiknow@gmail.com
Night	King	killerstare@gmail.com

WeaponOwner

Weapon	Name	Metal
Sword	Snow	Valyrian steel
Dagger	Lannister	Dragon glass

# Example Time

Person × Weapon (intermediate result)

FirstName	LastName	Email	Weapon	Name	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Snow	Valyrian steel
Jon	Snow	knowsnothing@hotmail.com	Dagger	Lannister	Dragon glass
Daenerys	Targaryen	bendtheknee@gmail.com	Sword	Snow	Valyrian steel
Daenerys	Targaryen	bendtheknee@gmail.com	Dagger	Lannister	Dragon glass
Tyrion	Lannister	idrinkandiknow@gmail.com	Sword	Snow	Valyrian steel
Tyrion	Lannister	idrinkandiknow@gmail.com	Dagger	Lannister	Dragon glass
Night	King	killerstare@gmail.com	Sword	Snow	Valyrian steel
Night	King	killerstare@gmail.com	Dagger	Lannister	Dragon glass

# Example Time

Person  $\bowtie$  Weapon  
LastName = Name

FirstName	LastName	Email	Weapon	Name	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Snow	Valyrian steel
Tyrion	Lannister	idrinkandiknow@gmail.com	Dagger	Lannister	Dragon glass

# Relational Algebra and SQL Statements





## The SELECT Statement: Detail

- A cut down version of the SELECT statement – MySQL
- **SELECT** [ALL | DISTINCT] *select\_expr* [, *select\_expr* ...]
  - List the columns (and expressions) that are returned from the query
- **FROM** *table\_references* ]
  - Indicate the table(s) or view(s) from where the data is obtained
- **WHERE** *where\_condition*
  - Indicate the conditions on whether a particular row will be in the result
- **GROUP BY** {*col\_name* | *expr*} [ASC | DESC], ...]
  - Indicate categorisation of results
- **HAVING** *where\_condition*
  - Indicate the conditions under which a particular category (group) is included in the result
- **ORDER BY** {*col\_name* | *expr* | *position*} [ASC | DESC], ...]
  - Sort the result based on the criteria
- **LIMIT** [{*offset*,} *row\_count* | *row\_count* OFFSET *offset*]
  - Limit which rows are returned by their return order (ie 5 rows, 5 rows from row 2)

Order is important! E.g. Limit cannot go before Group By or Having

# Relational Algebra and SQL Statements

Solutions [[Link to Jupyter Notebook](#)]

# Week 5 Lab

- Canvas → Modules → Week 5 → Lab → L05 SQL 1 (PDF)
- Objectives:
  - Install the lab schemas, tables and data
  - Learn SQL (Structured Query Language) SELECT syntax
  - Practise writing SQL queries
  - Join tables using natural and inner joins
- Breakout rooms, “ask for help” button if you need help or have any questions