

Database Systems

Tutorial Week 7

Objectives

- I. Effect of index on selection operator
- II. Matching index
- III. Cost estimation for different joins

Exercises (5 minutes)

1. Question about the effect of index on selection:

Consider a relation $R(a,b,c,d,e)$ containing 5,000,000 records, where each data page of the relation holds 10 records. R is organized as a sorted file with secondary indexes. Assume that $R.a$ is a candidate key for R , with values lying in the range 0 to 4,999,999, and that R is stored in $R.a$ order. For each of the following relational algebra queries, state which of the following three approaches is most likely to be the cheapest:

- Access the sorted file of R directly.
- Use a B+ tree index on attribute $R.a$.
- Use a hash index on attribute $R.a$.

Queries:

- $\sigma_{a < 50000}(R)$
- $\sigma_{a = 50000}(R)$
- $\sigma_{a > 50000 \wedge a < 50010}(R)$

Exercises

a. $\sigma_{a < 50000}(R)$

- Sorted file over R
- Start from the beginning of the file and stop when $a = 50,000$
- If we choose B+ tree index, we first have to find where $a = 50,000$ and then find records to the left (lower in value)

Exercises

b. $\sigma_{a=50000}(R)$

- Equality query
- Hash-index will be the most cost-effective

Exercises

c. $\sigma_{a > 50000 \wedge a < 50010}(R)$

- Range query that doesn't begin at the start of the sorted file
- B+ tree index will be the cheapest
- With sorted file, we'd have to do 50,000 comparisons before even getting to $a = 50,000$

Matching Predicates

- Queries usually have multiple conditions called **predicates**
- What's the predicate for Q2a?
 - `Sailors.sid < 50,000`
- A B-tree index matches predicate(s) that involve attributes in a **prefix of the search key**
- Say you have an index on `<a, b, c>`
- This index will match conditions/predicates on `(a, b, c)`, `(a, b)` and `(a)`, but not `(b)`, `(b, c)` or `(c)`
- Any combination of predicates that involve attributes in a **prefix** are called **matching predicates**
- An index can be used to speed up an analysis over matching predicates only

Exercises (10 minutes)

2. Matching index

Consider the following schema for the Sailors relation:

Sailors (sid INT, sname VARCHAR(50), rating INT, age DOUBLE)

For each of the following indexes, list whether the index matches the given selection conditions and briefly explain why.

- A B+ tree index on the search key (Sailors.sid)
 - a. $\sigma_{\text{Sailors.sid} < 50,000}(\text{Sailors})$
 - b. $\sigma_{\text{Sailors.sid} = 50,000}(\text{Sailors})$
- A hash index on the search key (Sailors.sid)
 - c. $\sigma_{\text{Sailors.sid} < 50,000}(\text{Sailors})$
 - d. $\sigma_{\text{Sailors.sid} = 50,000}(\text{Sailors})$
- A B+ tree index on the search key (Sailors.rating, Sailors.age)
 - e. $\sigma_{\text{Sailors.rating} < 8 \wedge \text{Sailors.age} = 21}(\text{Sailors})$
 - f. $\sigma_{\text{Sailors.rating} = 8}(\text{Sailors})$
 - g. $\sigma_{\text{Sailors.age} = 21}(\text{Sailors})$

Exercises

A B+ tree index on the search key (Sailors.sid)

- a. $\sigma_{\text{Sailors.sid} < 50,000}(\text{Sailors})$
- b. $\sigma_{\text{Sailors.sid} = 50,000}(\text{Sailors})$

With B+ tree indexes, we can do range checks and equality checks

- a. Yep! Matching predicates are $\text{Sailors.sid} < 50,000$
- b. Yep! Matching predicates are $\text{Sailors.sid} = 50,000$

Exercises

A hash index on the search key (Sailors.sid)

c. $\sigma_{\text{Sailors.sid} < 50,000}(\text{Sailors})$

d. $\sigma_{\text{Sailors.sid} = 50,000}(\text{Sailors})$

With hash indexes, we can do equality checks

c. Nope! Range queries can't be applied to a hash index :(

d. Yep! Matching predicates are $\text{Sailors.sid} = 50,000$

Exercises

A B+ tree index on the search key (Sailors.rating, Sailors.age)

- e. $\sigma_{\text{Sailors.rating} < 8 \wedge \text{Sailors.age} = 21}(\text{Sailors})$
- f. $\sigma_{\text{Sailors.rating} = 8}(\text{Sailors})$
- g. $\sigma_{\text{Sailors.age} = 21}(\text{Sailors})$

e. Yep! Matching predicates are $\text{Sailors.rating} < 8$, and
 $\text{Sailors.rating} < 8 \wedge \text{Sailors.age} = 21$

f. Yep! Matching predicates are $\text{Sailors.rating} = 8$

g. Nope! The index on (Sailors.rating, Sailors.age) is primarily sorted on Sailors.rating, so you'd need to search the entire relation to find sailors with a particular Sailors.age value

Exercises (15 mins)

3. Question about the cost analysis of different joins:

Consider the join $R \bowtie_{R.a = S.b} S$, given the following information about the relations to be joined:

- Relation R contains 10,000 tuples and has 10 tuples/page.
- Relation S contains 2,000 tuples and also has 10 tuples/page.
- Attribute b of relation S is the primary key for S.
- Both relations are stored as simple heap files.
- Neither relation has any indexes built on it.
- 52 buffer pages are available.

The cost metric is the number of page I/Os unless otherwise noted and the cost of writing out the result should be uniformly ignored.

- a. What is the cost of joining R and S using the **page-oriented Simple Nested Loops** algorithm? What is the minimum number of buffer pages (in memory) required in order for this cost to remain unchanged?
- b. What is the cost of joining R and S using the **Block Nested Loops** algorithm? What is the minimum number of buffer pages required in order for this cost to remain unchanged?
- c. What is the cost of joining R and S using the **Sort-Merge Join** algorithm? Assume that the external merge sort process can be completed in 2 passes.
- d. What is the cost of joining R and S using the **Hash Join** algorithm?
- e. What would the lowest possible I/O cost be for joining R and S using any join algorithm, and how much buffer space would be needed to achieve this cost? Explain briefly.

Exercises

Step 1: always find how many pages you need for each relation

Number of pages = number of tuples / number of tuples/page

Let M be the number of pages in R

- $M = 10,000 / 10 = 1,000$

Let N be the number of pages in S

- $N = 2,000 / 10 = 200$

Let B be the number of buffer pages available

- $B = 52$

Exercises

a. What is the cost of joining R and S using the **page-oriented Simple Nested Loops** algorithm?

- Need to know page-oriented Simple Nest Loop in detail
- Do a page-by-page scan of the outer relation
- For each outer page, do a page-by-page scan of the inner relation
- How do we minimise the cost of joining R and S?
 - Formula is: $\text{NPages}(\text{outer}) + \text{NPages}(\text{outer}) \times \text{NPages}(\text{inner})$
 - Select the smaller relation as the outer relation
 - S has 200 pages, R has 1,000 pages, so in this case, choose S as the outer relation

Exercises

a. What is the cost of joining R and S using the **page-oriented Simple Nested Loops** algorithm?

- $\text{Cost} = \text{NPages}(\text{outer}) + \text{NPages}(\text{outer}) \times \text{NPages}(\text{inner})$
- $S = \text{outer}, R = \text{inner}$
- Cost
 - $= 200 + (200 \times 1,000)$
 - $= 200,200 \text{ I/O}$

Exercises

a. What is the minimum number of buffer pages (in memory) required in order for this cost to remain unchanged?

- In the page-oriented Simple Nested Loops algorithm, we don't use multiple buffers at a time
- Minimum number of buffer pages is:
 - One buffer page for the left input
 - One buffer page for the right input
 - One buffer page to store the output/result
 - Total = 3 buffer pages required

Exercises

b. What is the cost of joining R and S using the **Block Nested Loops** algorithm?

- Need to know Block Nested Loop in detail
- Outer relation is read in “blocks”
 - “Blocks” are groups of pages that will fit into whatever buffer pages are available
- For each block, do a page-by-page scan of the inner relation
- Each page of the outer relation is scanned once
- Each page of the inner relation is scanned once per block
- $\text{Cost} = \text{NPages}(\text{outer}) + \text{NBlocks} \times \text{NPages}(\text{inner})$
- Use one buffer page for scanning the inner relation
- Use one buffer page to store the output
- Use all other buffer pages to hold the blocks of the outer relation
- $\text{NBlocks} = \text{ceil}(\text{NPages}(\text{outer}) / B - 2)$
 - Ceiling means you round up to the nearest integer

Exercises

b. What is the cost of joining R and S using the **Block Nested Loops** algorithm?

- Cost
 - = $\text{NPages}(\text{outer}) + \text{ceil}(\text{NPages}(\text{outer}) / B - 2) \times \text{NPages}(\text{inner})$
 - = $200 + \text{ceil}(200/50) \times 1,000$
 - = $200 + (4 \times 1,000)$
 - = 4,200 I/O

Exercises

b. What is the minimum number of buffer pages required in order for this cost to remain unchanged?

- $\text{Cost} = \text{NPages}(\text{outer}) + \text{NBlocks} \times \text{NPages}(\text{inner})$
- $\text{NBlocks} = \text{ceil}(\text{NPages}(\text{outer}) / B - 2)$
- Fewer buffer pages \rightarrow denominator decreases \rightarrow NBlocks increases \rightarrow overall cost increases
- So we can't use fewer buffer pages
- Have to use all 52 i.e. minimum = 52

Exercises

c. What is the cost of joining R and S using the Sort-Merge Join algorithm?
Assume that the external merge sort process can be completed in **2 passes**.

- Don't need to know Sort-Merge Join algorithm in too much detail
- Cost

$$\begin{aligned} &= \text{NPages(outer)} + \text{NPages(inner)} && \leftarrow \text{Cost of merging outer and inner relations} \\ &+ 2 \times \text{NPages(outer)} \times \text{num_passes(outer)} && \leftarrow \text{Cost of accessing and sorting outer relation} \\ &+ 2 \times \text{NPages(inner)} \times \text{num_passes(inner)} && \leftarrow \text{Cost of accessing and sorting inner relation} \end{aligned}$$

$$= 200 + 1000 + (2 \times 200 \times 2) + (2 \times 1,000 \times 2)$$

$$= 6,000 \text{ I/O}$$

Exercises

d. What is the cost of joining R and S using the Hash Join algorithm?

- Don't need to know Hash Join algorithm in too much detail
- In hash join, each relation is partitioned and then the join is performed by “matching” elements from corresponding partitions
- Cost
 - = $3 \times (\text{NPages}(\text{outer}) + \text{NPages}(\text{inner}))$
 - = $3 \times (200 + 1,000)$
 - = 3,600 I/O

Exercises

e. What would the lowest possible I/O cost be for joining R and S using any join algorithm, and how much buffer space would be needed to achieve this cost?

- Optimal cost achieved if each relation is read only once
- We could store entire smaller relation in memory
- We could then read in the larger relation page by page
 - For each tuple in the larger relation, we could search the smaller relation (which exists entirely in memory) for matching tuples
- Total cost = $\text{NPages}(\text{smaller relation}) + \text{NPages}(\text{bigger relation}) = 200 + 1,000 = 1,200$ I/O
- Buffer would have to:
 - Hold the smaller relation
 - Have one buffer page to read in the larger relation
 - Have one buffer page to store the output
- Minimum number of buffer pages is $\text{NPages}(\text{smaller relation}) + 1 + 1 = 200 + 1 + 1 = 202$

Week 7 Lab

- Canvas → Modules → Week 7 → Lab → L07 SQL 3 (PDF)
- Objectives:
 - Practice further joins involving three and four tables
 - Understand CASE statements and the UNION clause
 - Develop complex SQL queries using derived tables and views
 - Note you're **not** allowed to use views for assignment 2
 - Create and understand relational divides using EXISTS and NOT EXISTS
 - Video on Canvas
- Breakout rooms, “ask for help” button if you need help or have any questions